

Management Sciences Topics: Convex Optimization

Final Project

1 Problem setup

We need to solve the optimization problem of a one-hidden-layer neural network

$$\min_{x_k \in \mathbb{R}^d, y_k \in \mathbb{R}, z \in \mathbb{R}^K, w \in \mathbb{R}} \frac{1}{n} \sum_{i=1}^n \mathcal{L} \left(b_i w + b_i \sum_{k=1}^K \sigma(a_i^\top x_k + y_k) z_k \right), \quad (1.1)$$

where K is the number of neurons, $a_i \in \mathbb{R}^d$ is a data point, $b_i \in \{-1, 1\}$ is the class label of a_i , $\sigma(z) = \max(z, 0)$ or $\frac{\exp(z)}{1+\exp(z)}$, $\mathcal{L}(z) = \max(1 - z, 0)$ or $\log(1 + \exp(-z))$.

2 Stochastic subgradient method

We first consider the subgradient with respect to each variables.

First, define the variables

$$X = [x_1, x_2, \dots, x_K] \in \mathbb{R}^{d \times K}, \quad Y = [y_1, y_2, \dots, y_K]^\top \in \mathbb{R}^{K \times 1}, \quad Z = [z_1, z_2, \dots, z_K]^\top \in \mathbb{R}^{K \times 1}. \quad (2.1)$$

Then a forward pass through the network can be written as

$$\begin{aligned} A_1 &= AX \oplus Y^\top, \\ A_2 &= \sigma(A_1)Z \oplus w \\ f &= \frac{1}{n} 1^\top L(b \odot A_2). \end{aligned} \quad (2.2)$$

By chain rule, the subgradients of f with respect to each variable are

$$\begin{aligned} \frac{\partial f}{\partial A_2} &= \frac{1}{n} L'(b \odot A_2) \odot b \\ \frac{\partial f}{\partial w} &= \frac{\partial f}{\partial A_2} \frac{\partial A_2}{\partial w} = \text{rowsum} \left(\frac{\partial f}{\partial A_2} \odot 1 \right) = \left(\frac{\partial f}{\partial A_2} \right)^\top 1, \\ \frac{\partial f}{\partial Z} &= \frac{\partial f}{\partial A_2} \frac{\partial A_2}{\partial Z} = \sigma(A_1)^\top \frac{\partial f}{\partial A_2}, \\ \frac{\partial f}{\partial A_1} &= \frac{\partial f}{\partial A_2} \frac{\partial A_2}{\partial A_1} = \frac{\partial f}{\partial A_2} Z^\top \odot \sigma'(A_1), \\ \frac{\partial f}{\partial Y} &= \frac{\partial f}{\partial A_1} \frac{\partial A_1}{\partial Y} = \left(\frac{\partial f}{\partial A_1} \right)^\top 1, \\ \frac{\partial f}{\partial X} &= \frac{\partial f}{\partial A_1} \frac{\partial A_1}{\partial X} = A^\top \frac{\partial f}{\partial A_1}. \end{aligned} \quad (2.3)$$

The stochastic subgradients can be chosen to be the subgradient when input is a minibatch of the whole dataset, i.e.

$$G(x, \xi_i) = \partial_x f(x; A_{\xi_i}, b_{\xi_i}) \quad (2.4)$$

for each variable x , where ξ_i is a uniformly sample index set for each i . We can control the size of each ξ_i to vary from online learning to full-batch learning.

3 Accelerated proximal gradient method

In order to use APG for this problem, we need to choose

$$\sigma(z) = \frac{\exp(z)}{1 + \exp(z)}, \quad \mathcal{L}(z) = \log(1 + \exp(-z)) \quad (3.1)$$

to guarantee the objective function is smooth. Notice

$$\mathcal{L}'(z) = -\frac{1}{1 + \exp(z)}, \quad (3.2)$$

and

$$|\mathcal{L}''(z)| = \left| \frac{e^z}{(1 + e^z)^2} \right| \leq \frac{1}{4}, \quad (3.3)$$

by Lagrange mean value theorem, we know the Lipschitz constant for \mathcal{L}' is $L = \frac{1}{4}$.

Now let's consider the Lipschitz constant for each derivatives.

For ∂_{A_2} ,

$$\begin{aligned} \left| \frac{\partial f}{\partial A_2^1} - \frac{\partial f}{\partial A_2^2} \right| &= \frac{1}{n} |\mathcal{L}'(b \odot A_2^1) \odot b - \mathcal{L}'(b \odot A_2^2) \odot b| = \frac{1}{n} |(\mathcal{L}'(b \odot A_2^1) - \mathcal{L}'(b \odot A_2^2)) \odot b| \\ &= \frac{1}{n} |\mathcal{L}''(\xi)(b \odot (A_2^1 - A_2^2)) \odot b| \\ &\leq \frac{1}{4n} |A_2^1 - A_2^2|. \end{aligned} \quad (3.4)$$

Then the Lipschitz constant for $\frac{\partial f}{\partial A_2}$ is $L_{A_2} = \frac{1}{4n}$.

For $\frac{\partial f}{\partial w}$,

$$\begin{aligned} \left| \frac{\partial f}{\partial w^1} - \frac{\partial f}{\partial w^2} \right| &= \left(\frac{\partial f}{\partial A_2^1} - \frac{\partial f}{\partial A_2^2} \right)^\top 1 \\ &\leq \frac{1}{4n} |w_1 - w_2| n = \frac{1}{4} |w_1 - w_2|. \end{aligned} \quad (3.5)$$

Then the Lipschitz constant for ∂_w is $L_w = \frac{1}{4}$.

For $\frac{\partial f}{\partial Z}$, we first notice $\sigma(x) \in (0, 1)$. Then

$$\begin{aligned} \left| \frac{\partial f}{\partial Z^1} - \frac{\partial f}{\partial Z^2} \right| &= \left| \sigma(A_1)^\top \left(\frac{\partial f}{\partial A_2^1} - \frac{\partial f}{\partial A_2^2} \right) \right| \leq \frac{1}{4} |\sigma(A_1)^\top (A_2^1 - A_2^2)| \\ &= \frac{1}{4} |\sigma(A_1)^\top \sigma(A_1) (Z_1 - Z_2)| \\ &\leq \frac{1}{4} \|\sigma(A_1)^\top \sigma(A_1)\|_2 |Z_1 - Z_2| \\ &\leq \frac{1}{4} \|\sigma(A_1)^\top \sigma(A_1)\|_F |Z_1 - Z_2| \\ &\leq \frac{1}{4} \sqrt{k^2 n^4} |Z_1 - Z_2| = \frac{1}{4} k n^2 |Z_1 - Z_2|. \end{aligned} \quad (3.6)$$

Then $L_z = \frac{1}{4} k n^2$ can be an upper bound for the Lipschitz constant for ∂_z .

The lipschitz constant for $\frac{\partial f}{\partial y}$ and $\frac{\partial f}{\partial x}$ are too complex to solve.

4 Proximal Gradient method with line search

In order to use PG for this problem, we need to choose

$$\sigma(z) = \frac{\exp(z)}{1 + \exp(z)}, \quad \mathcal{L}(z) = \log(1 + \exp(-z)) \quad (4.1)$$

to guarantee the objective function is smooth.

Since we use an iterative way to update the variables, the value of x will be updated from x_0 to x_1 before we start to update y, z, w . So there are two schemes in updating:

1. Use x_0 to update y , and use x_0, y_0 to update z , and use x_0, y_0, z_0 to update w (referred to as PG1).
2. Use x_1 to update y . In this case, we need to recompute the objective function and the derivatives using x_1 . For y, z, w , we also use this scheme (referred to as PG2).

We test both schemes in the experiments.

5 Experiments

For the experiments, we choose the number of hidden layers $K = 128$, and

$$\sigma(z) = \frac{\exp(z)}{1 + \exp(z)}, \quad \mathcal{L}(z) = \log(1 + \exp(-z)). \quad (5.1)$$

All variables x, y, z, w are initialized from a uniform distribution on $(-1, 1)$.

We test this one-hidden-layer neural network on the datasets `rcv1.binary` and `covtype` from libsvm library. We use the same random seed across the experiments to get the same random initialization for a fair comparison. For all following experiments, we compute the objective value, subgradient and out-of-sample accuracy on the averaged variable after each iteration.

We compute the out-of-sample prediction accuracy by

$$Acc = \frac{1}{n} \sum_{i=1}^n 1_{\text{sign}(A_{2i}) = \text{sign}(b_i)} \times 100\% \quad (5.2)$$

on an independent test set.

We compute 2-norm of ∂x by first flatten it to a vector, and then compute the vector 2-norm instead of matrix 2-norm.

5.1 rcv1.binary

For `rcv1.binary`, we use only the `rcv1.train.binary` from libsvm library. We randomly sample 20% of the dataset, i.e., 4048 samples to form the test set, and use the rest part as the training set. Both test set and training set remain the same across the four methods.

For memory issues (i.e., since we need to keep track of each variables and their derivatives), for PG and APG, we run the algorithm with $K = 200$ iterations and only keep track of the results from $2^{nd}, 4^{th}, \dots, 200^{th}$ iterations. For SSG, we run $K = 1200$ iterations with minibatch size $N = 2699$ in each iteration, and only keep track of the results from $12^{th}, 24^{th}, \dots, 1200^{th}$ iterations. Then the total numbers of training samples passed for all three algorithms are the same.

The objective and accuracy of the four methods are shown in Figure 5.1. PG2 achieves the lowest objective and highest accuracy among the four methods, with PG1 follows. However, for PG with line search, the computational complexity is much higher than APG and SSG because of the inner loop, and the computational cost of PG2 is much higher than PG1. Figure 5.2 shows the 2-norm of each subgradient with respect to the number of iterations. The result agrees with the result of objective and accuracy.




Figure 5.1: Objective and Accuracy of the four methods with respect to the number of iterations.




Figure 5.2: Norm of subgradients of the four methods with respect to the number of iterations.

5.2 covtype.binary

We also test the four methods on the `covtype.binary`. We first rescale the b vector to $\{-1, 1\}$ as mentioned in a homework before. We randomly sampled 20% of the whole dataset, i.e., 116202 samples to form the test set, and use the rest part as the training set.

For APG and PG, we run 200 iterations and record the results from the even iterations. For SSG, we run 2000 iterations with minibatch size $N = 46481$, and only record the results from the 20^{th} , 40^{th} , \dots , 2000^{th} iterations.

5.3 Bernoulli distribution

For this part, we generate 10000 Gaussian-distributed samples of dimension $d = 4$ as the dataset A , and sample each variable x, y, z, w from the uniform distribution on $(-1, 1)$.

We train a 128-hidden-unit neural network on this dataset. SSG is run $K = 2000$ iterations with minibatch size $N = 1000$, and PG, APG are run $K = 200$ iterations.

The objective value on each iteration and the prediction accuracy on the training set is shown in Figure 5.3. The relative error of each variable x, y, z, w measured in relative L_2 norm is shown in Figure 5.4. We can observe that PG2 and SSG achieve the fastest convergence in both accuracy and objective value. However, a clear overfitting occurs for almost all variables, so we cannot decide which method has the smallest error in the variables.




Figure 5.3: Objective and Accuracy of the four methods with respect to the number of iterations.




Figure 5.4: Relative L_2 error for each variable with respect to the number of iterations.